

# Transactions, Constraints, Logic Code in Database

CSC365

# Command Syntax

## START TRANSACTION

*[transaction\_characteristic [, transaction\_characteristic] ...]*

*transaction\_characteristic*: {

WITH CONSISTENT SNAPSHOT

| READ WRITE

| READ ONLY

}

BEGIN [WORK]

COMMIT [WORK] [AND [NO] CHAIN] [[NO] RELEASE]

ROLLBACK [WORK] [AND [NO] CHAIN] [[NO] RELEASE]

SET autocommit = {0 | 1}

# Example

```
START TRANSACTION;
```

```
SELECT @A:=SUM(salary) FROM table1 WHERE type=1;
```

```
UPDATE table2 SET summary=@A WHERE type=1;
```

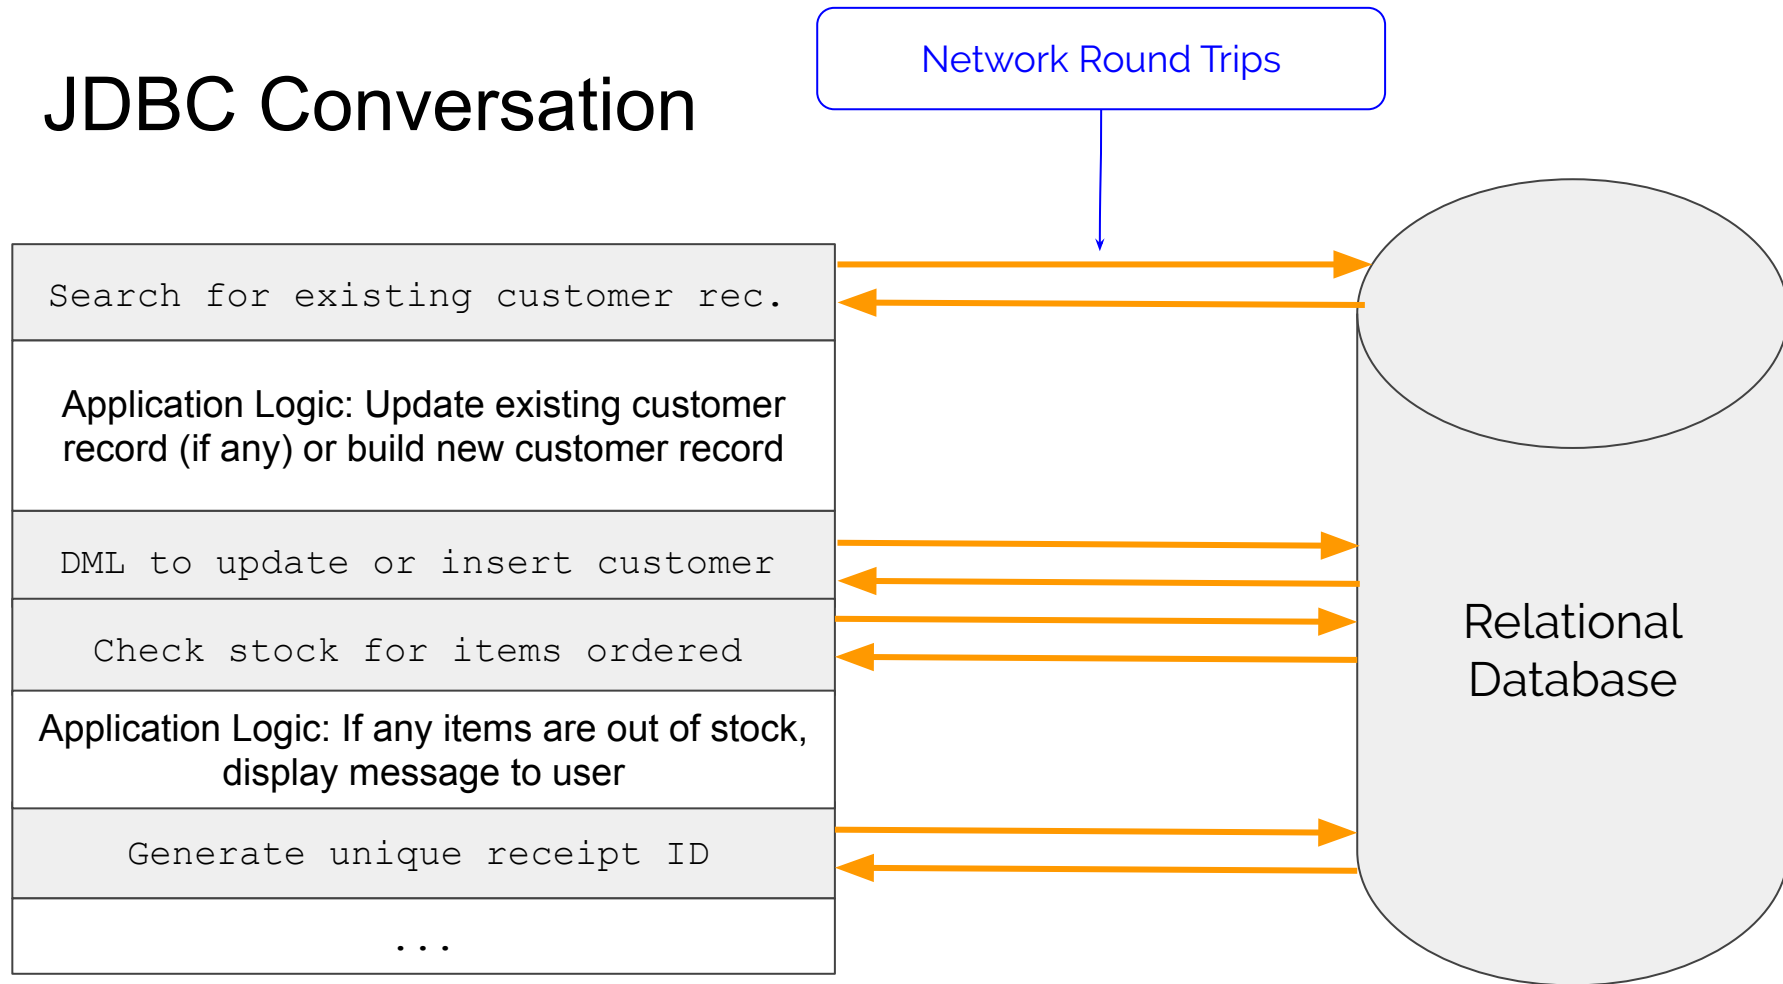
```
COMMIT;
```

# JDBC Applications

A JDBC application typically interacts with a database in a *conversational* way: executing multiple SQL statements, each involving a round trip to the database. Consider a hypothetical customer order:

1. Check whether a customer record exists
2. Update existing customer or create a new customer record
3. Confirm stock level of ordered products
4. Save order detail (receipt & line items)
5. Record payment detail

# JDBC Conversation



# Multi-Step Conversation: Approaches

- Regular JDBC Statement for each step in process
- Use PreparedStatement to avoid cost of repeated query parse/compile in Java
- Use batch mode to group multiple SQL statements
- Stored Procedures

# JDBC - Batch Processing

- The `addBatch()` method of *Statement*, *PreparedStatement*, and *CallableStatement* is used to add individual statements to the batch. The `executeBatch()` is used to start the execution of all the statements grouped together.
- `executeBatch()` - returns an array of integers, and each element of the array represents the update count for the respective update statement.
- `clearBatch()` - This method removes all the statements you added with the `addBatch()` method. However, you cannot selectively choose which statement to remove.

# JDBC Batch Example

```
// Create SQL statement
String SQL = "INSERT INTO Employees (id, first, last, age) VALUES (?, ?, ?, ?)";
// Create PreparedStatement object
PreparedStatement pstmt = conn.prepareStatement(SQL);

//Set auto-commit to false
conn.setAutoCommit(false);

// Set the variables
pstmt.setInt( 1, 400 );
pstmt.setString( 2, "Pappu" );
pstmt.setString( 3, "Singh" );
pstmt.setInt( 4, 33 );
// Add it to the batch
pstmt.addBatch();

//add more batches repeating the Set the Variables and the Add it to the Batch steps
//Create an int[] to hold returned values
int[] count = stmt.executeBatch();
//Explicitly commit statements to apply changes
conn.commit();
```



# Stored Procedures

The SQL standard specifies support for *Persistent Stored Modules* (PSM), commonly referred to as Stored Procedures. With this support, procedural code can be stored in the database and executed on the database server. Although the SQL standard defines a standard programming language, vendor implementations are splintered:

- Oracle defines [PL/SQL](#)
- [MySQL supports a variant of PL/SQL](#)
- Microsoft SQL Server implements Transact-SQL, often abbreviated [T-SQL](#)

# Stored Procedures - Pros & Cons

Consider "plain" JDBC/SQL versus stored procedures.

Let's list pros & cons of stored procedures.

# Stored Procedures - Advantages

- Possible performance gains:
  - Fewer round trips during multi-step conversations
  - Cached / precompiled SQL statements
- Shields applications from details about data structure
  - Security
  - Easier schema evolution
- Common logic which may be shared by multiple applications; no need to re-implement across multiple applications that may share similar needs

# Stored Procedures - Disadvantages

- Syntax is often an odd mix of Pascal/Ada and SQL
  - No standardization, every current RDBMS defines its own language
  - Developers must learn specifics
  - Limited extendibility (ie. no "import" capability)
- Debugging and testing are difficult when logic spans both the application and database
- Release / versioning picture is hazy
  - DBAs often reluctant to apply frequent updates due to unknown impact on DB security and stability

# Trigger

A trigger is a named database object that is associated with a table, and that activates when a particular event occurs for the table.

The trigger becomes associated with the table named ***tbl\_name***, which must refer to a permanent table.

Triggers are typically written in the procedural language supported by the RDBMS (ie. PL/SQL, T-SQL)

# Trigger Events

At a basic level, a trigger may be defined to execute BEFORE or AFTER any of these events:

- INSERT
- UPDATE
- DELETE
- ~~SELECT~~

Each trigger is attached to a single table.

# DELETE / UPDATE Triggers

In a DELETE trigger, the OLD keyword allows you to access columns in the row(s) affected by the delete.

UPDATE triggers allow you to reference OLD and NEW:

- OLD refers to the row prior to update (OLD is read-only)
- NEW enables you to access (and possibly change) the updated row values

# Example Trigger Use Cases

- Enforce assertions / complex constraints
  - Recall: assertions are constraints that cannot readily be expressed using basic relational model integrity constraints (PK, FK, nullability, CHECK)
- Maintain database integrity
  - Create mandatory "child" records
- Audit logging



# Trigger Considerations

Triggers and stored procedures share many of the same pros & cons. Some concerns specific to triggers:

- Performance: trigger code must be extremely efficient, since it runs for *every* record added/updated/removed via DML
- Sequencing of multiple triggers on the same table / same event(s) (if supported)
  - MySQL allows you to specify PRECEDES or FOLLOWS in trigger definition

# Using Triggers to Enforce Constraints

It is possible to use BEFORE triggers to prevent an INSERT/UPDATE/DELETE from taking place.

This capability can be used to enforce constraints that go beyond the basic constraints that may be specified in SQL DDL.

# Triggers

CREATE

[DEFINER = *user*]

TRIGGER *trigger\_name*

*trigger\_time* *trigger\_event*

ON *tbl\_name* FOR EACH ROW

[*trigger\_order*]

*trigger\_body*

*trigger\_time*: { BEFORE | AFTER }

*trigger\_event*: { INSERT | UPDATE | DELETE }

*trigger\_order*: { FOLLOWS | PRECEDES } *other\_trigger\_name*

# Trigger Examples

```
-- before insert
DELIMITER $
CREATE TRIGGER `creditcard_before_insert` BEFORE INSERT ON
`CreditCard`
FOR EACH ROW
BEGIN
    IF NEW.balance > NEW.limit THEN
        SIGNAL SQLSTATE '99999'
        SET MESSAGE_TEXT = 'check constraint on CreditCard.balance failed';

    END IF;
END$
DELIMITER ;
```

# Trigger Examples

```
-- before update
DELIMITER $
CREATE TRIGGER `creditcard_before_update` BEFORE UPDATE ON
`CreditCard`
FOR EACH ROW
BEGIN
    IF NEW.balance > OLD.limit THEN
        SIGNAL SQLSTATE '99999'
        SET MESSAGE_TEXT = 'check constraint on CreditCard.balance failed';

    END IF;
END$
DELIMITER ;
```

# Trigger Examples

```
-- before insert
DELIMITER $
CREATE TRIGGER `customer_before_insert` BEFORE INSERT ON `Customer`
FOR EACH ROW
BEGIN
    IF NEW.ssn < 100000000 THEN
        SIGNAL SQLSTATE '12345'
        SET MESSAGE_TEXT = 'check constraint on Customer.ssn failed';

    END IF;
END$
DELIMITER ;
```

# Trigger Examples

```
-- before update
DELIMITER $
CREATE TRIGGER `customer_before_update` BEFORE UPDATE ON `Customer`
FOR EACH ROW
BEGIN
    IF NEW.ssn < 100000000 THEN
        SIGNAL SQLSTATE '12345'
        SET MESSAGE_TEXT = 'check constraint on Customer.ssn failed';

    END IF;
END$
DELIMITER ;
```

# Stored Procedure Syntax

## CREATE

```
[DEFINER = user]  
PROCEDURE sp_name ([proc_parameter[,...]])  
[characteristic ...] routine_body
```

## CREATE

```
[DEFINER = user]  
FUNCTION sp_name ([func_parameter[,...]])  
RETURNS type  
[characteristic ...] routine_body
```

*proc\_parameter:*

```
[ IN | OUT | INOUT ] param_name type
```

*func\_parameter:*

```
param_name type
```

*type:*

*Any valid MySQL data type*

*characteristic:*

```
COMMENT 'string'
```

```
| LANGUAGE SQL
```

```
| [NOT] DETERMINISTIC
```

```
| { CONTAINS SQL | NO SQL | READS SQL DATA |
```

```
MODIFIES SQL DATA }
```

```
| SQL SECURITY { DEFINER | INVOKER }
```

*routine\_body:*

*Valid SQL routine statement*



# Stored Procedure Example

```
DELIMITER $
CREATE PROCEDURE `check_balance`(IN balance DECIMAL(10,2), IN lim
DECIMAL(10,2))
BEGIN
    IF balance > lim THEN
        SIGNAL SQLSTATE '99999'
        SET MESSAGE_TEXT = 'check constraint on CreditCard.balance failed';
    END IF;
END$
DELIMITER ;
```

# Stored Procedure Example

```
-- before update
DELIMITER $
CREATE TRIGGER `creditcard_before_update` BEFORE UPDATE ON
`CreditCard`
FOR EACH ROW
BEGIN
    CALL check_balance(NEW.balance, OLD.limit);
END$
DELIMITER ;
```

# Functions

In addition to stored procedures, many databases support user-defined **functions**. Functions are stored in the database catalog, again using an RDBMS-specific language such as PL/SQL.

Unlike stored procedures, functions must return a value and may be used within `SELECT` statements.

# User-Defined Function Example

```
-- DROP FUNCTION miles_to_km;  
-- No ; chars in the function body, no need to change delimiter  
CREATE FUNCTION miles_to_km (miles DECIMAL(15,5))  
RETURNS DECIMAL(15,5) DETERMINISTIC  
    RETURN miles * 1.60934;
```

```
SELECT miles_to_km(100.0);
```

DETERMINISTIC tells MySQL that the function always produces the same result for the same input

# Dropping Triggers, Stored Procedures

```
DROP TRIGGER `trigger_name`;
```

```
DROP PROCEDURE `procedure_name`;
```

```
DROP FUNCTION `function_name`;
```

# Constraints

# Expressing Constraints in Relational Algebra

- $R = \emptyset$ 
  - “The value of R must be empty.”
- $R \subseteq S$ 
  - “Every tuple in the result of R must also be in the result of S.”

# Algebraic Primary Key Constraint

The primary key (PK) must be *unique* across all tuples. Following from this, if two tuples agree on the PK, they must also agree on *all* other attributes.

Expressed as an algebraic constraint:

$$\sigma_{R1.<PK> = R2.<PK> \text{ AND } R1.<\text{non key attribute}> \neq R2.<\text{non key attribute}>} (\varrho_{R1}(R) \times \varrho_{R2}(R)) = \emptyset$$



# Referential Integrity Constraints

A foreign key (FK) value must also be a key value of a relation in the same database. Expressing this as an algebraic constraint using set containment:  $\pi_{FK}(R) \subseteq \pi_K(S)$

or, equivalently:

$$\pi_{FK}(R) - \pi_K(S) = \emptyset$$

$$\pi_{owner\_id}(\text{Ownership}) \subseteq \pi_{id}(\text{Customer}) \text{ or}$$
$$\pi_{owner\_id}(\text{Ownership}) - \pi_{id}(\text{Customer}) = \emptyset$$

# Expressing Other Constraints

The balance on a Credit Card must be less than its credit limit

$$\sigma_{R.limit < R.balance}(\varrho_R(\text{CreditCard})) = \emptyset$$

# Setting Constraints in DDL (SQL)

```
CREATE TABLE CreditCard (  
    number INT NOT NULL AUTO_INCREMENT,  
    type ENUM('Visa', 'Master Card', 'American Express', 'Discover') NOT NULL,  
    `limit` INT,  
    balance INT,  
    active TINYINT(1),  
    PRIMARY KEY (number),  
    CHECK (`limit` >= balance)  
);
```

CHECK is now supported by MySQL as of v8.0.  
Use triggers in older versions.