

Writing Queries involving more than one Relation

CSC365
Spring 2019

Products and Joins in SQL

- To couple relations in one query
 - List relations separated by comma in FROM clause
 - `SELECT * FROM Movies, Genre;`

Disambiguating Attributes

- `SELECT mid, title, year, genre FROM Movies, Genre WHERE mid=mid;`
 - The column name `mid` is ambiguous.
- To disambiguate, we can prepend table name to column names.

```
SELECT Movies.mid, title, year, genre
```

```
FROM Movies, Genre
```

```
WHERE Movies.mid=Genre.mid;
```

Tuple Variables (Renaming Relations ρ)

- Sometimes, you want to alias the name of a relation to a short string

```
SELECT m.mid, title, year, genre
```

```
FROM Movies as m, Genre as g
```

```
WHERE m.mid=g.mid;
```

Eliminating Duplicates

DISTINCT

```
SELECT DISTINCT m.mid, title, year, genre  
FROM Movies as m, Genre as g  
WHERE m.mid=g.mid;
```

Cartesian Product

$R \times S$

There are two ways to express cartesian product in SQL.

Implicitly:

```
SELECT *  
FROM Movies, Genre;
```

Explicitly:

```
SELECT *  
FROM Movies CROSS JOIN Genre;
```

Theta Joins

We can use the following simple form to express theta joins:

Relational algebra expression:

$$U \bowtie_{A < V.C \text{ AND } U.B \neq V.B} V$$

```
SELECT *  
FROM U, V  
WHERE A < V.C AND U.B <> V.B
```

Theta Joins

```
SELECT *
```

```
FROM Movies m, Genre g
```

```
WHERE m.mid = g.mid AND g.genre = 'Sci-Fi';
```


Equi Joins

The SQL standard defines JOIN ... USING as a shortcut for natural equijoins. This allows explicit control over the join columns (vs. NATURAL JOIN)

```
SELECT *
```

```
FROM Movies INNER JOIN Genre USING (mid);
```

This syntax requires columns with the same name in joined tables.

Natural Joins

There are multiple ways to express a natural join in SQL.

Explicit projection and equijoin on common column(s):

```
SELECT DISTINCT Movies.*, Genre.genre  
FROM Movies, Genre  
WHERE Movies.mid = Genre.mid;
```

Implicitly, using NATURAL JOIN:

```
SELECT DISTINCT *  
FROM Movies NATURAL JOIN Genre;
```

Theta Joins with JOIN

As an alternative, we can use the following to express theta joins. Note the lack of a WHERE clause. Instead, we use the ANSI ***infix*** operator, [INNER] JOIN ... ON, to specify our join condition. The INNER keyword is *optional*.

```
SELECT *  
FROM U INNER JOIN V ON A < V.C AND U.B <> V.B
```

Relational algebra expression (same as previous slide):

$$U \bowtie_{A < V.C \text{ AND } U.B \neq V.B} V$$

Theta Joins with JOIN

```
SELECT m.*, g.genre
```

```
FROM Movies m
```

```
JOIN Genre g ON m.mid = g.mid AND g.genre = 'Sci-Fi';
```

SQL SELECT - Theta Join - WHERE vs JOIN ... ON

Two ways to express theta join:

```
SELECT * FROM U, V WHERE A < V.C AND U.B <> V.B
```

```
SELECT * FROM U INNER JOIN V ON A < V.C AND U.B <> V.B
```

- What's the difference?
 - JOIN ON is the ANSI standard
 - Readability, especially with more than one join
 - Other types of joins we'll see in SQL (OUTER)

Recommendation: Use JOIN ... ON for joins, WHERE clause for non-join conditions

Multiple Joins

A single SELECT statement may include any number of JOINS:

```
SELECT m.*, g.genre  
FROM Movies m  
JOIN Genre g ON g.mid = m.mid AND g.genre = 'Sci-Fi'  
JOIN StarsIn s ON s.mid = m.mid AND sname = 'Will Smith';
```

For inner joins, the join order has no impact on results.

JOINS & WHERE

A SELECT statement may combine JOIN and WHERE:

```
SELECT m.*, g.genre  
FROM Movies m  
JOIN Genre g ON g.mid = m.mid AND g.genre = 'Sci-Fi'  
WHERE m.year < 2000;
```

Corresponding Relational Algebra expression tree?

Aggregation Operators

SQL SELECT also supports **Aggregate Functions**, which summarize multiple rows.

- COUNT()
- MIN()
- MAX()
- SUM()
- AVG()

COUNT()

```
SELECT COUNT(*) FROM Movies;
```

COUNT(*)
482

```
SELECT COUNT(*)
```

```
FROM Movies
```

```
WHERE director = 'Steven Spielberg';
```

COUNT(*)
6

MIN()

```
SELECT MIN(imdb) FROM Movies;
```

```
SELECT MIN(imdb) FROM Movies WHERE director = 'Steven Spielberg';
```

```
SELECT MIN(title) FROM Movies WHERE director = 'Steven Spielberg';
```

+	-----	+
	MIN(title)	
+	-----	+
	A.I. Artificial Intelligence	
+	-----	+

MAX()

```
SELECT MAX(imdb) FROM Movies;
```

```
SELECT MAX(imdb) FROM Movies WHERE director = 'Steven Spielberg';
```

```
SELECT MAX(title) FROM Movies WHERE director = 'Steven Spielberg';
```

MAX(title)
War of the Worlds

SUM()

SELECT SUM(gross)/1000000 as grossInMillion

FROM Movies

WHERE director = 'Steven Spielberg';

title	grossInMillion
The BFG	52.7923
The Adventures of Tintin	77.5640
A.I. Artificial Intelligence	78.6167
Minority Report	132.0141
War of the Worlds	234.2771
Indiana Jones and the Kingdom of the Crystal Skull	317.0111

grossInMillion
892.2753

AVG()

```
SELECT AVG(imdb) FROM Movies;
```

```
SELECT AVG(imdb) FROM Movies WHERE director = 'Steven Spielberg';
```

```
SELECT AVG(gross)/1000000 as grossInMillion FROM Movies;
```

```
SELECT AVG(gross)/1000000 as grossInMillion
```

```
FROM Movies
```

```
WHERE director = 'Steven Spielberg';
```

Aggregate Functions: ALL vs DISTINCT

By default, SQL aggregate functions operate in "ALL" mode:

```
SELECT COUNT(sname), COUNT(ALL sname) FROM StarsIn;
```

DISTINCT causes the aggregate function to consider only unique values:

```
SELECT COUNT(sname), COUNT(DISTINCT sname) FROM StarsIn;
```

COUNT(sname)	COUNT(ALL sname)
1446	1446

COUNT(sname)	COUNT(DISTINCT sname)
1446	811

The functions AVG(), COUNT(), and SUM() support DISTINCT mode

SQL SELECT - Aggregate Functions

- Used alone, an aggregate function collapses all rows into *one* summarized row.
- What if we want to include other column values in our result set?

```
SELECT director, AVG(imdb)  
FROM Movies;
```

Grouping - GROUP BY

The GROUP BY clause, along with aggregate functions, allows us to identify which column(s) form the basis for our summary:

```
SELECT director, AVG(imdb)
```

```
FROM Movies
```

```
GROUP BY director;
```


GROUP BY

GROUP BY is permitted on both columns and the result of *scalar* functions:

```
SELECT
    FLOOR(DATEDIFF(CURRENT_DATE, birthdate) / 365) as age,
    COUNT(birthdate)
FROM Stars
GROUP BY FLOOR(DATEDIFF(CURRENT_DATE, birthdate) / 365);
```

GROUP BY with Conditions

Conditions on aggregate values require special treatment.

List all directors and average gross in million dollars per director of Movies made in or after year 2000 sorted in the order of the average gross.

```
SELECT director, AVG(gross)/1000000 as grossInMillion  
FROM Movies  
WHERE year >= 2000  
GROUP BY director  
ORDER BY grossInMillion;
```

GROUP BY with Conditions

The HAVING clause, used along with GROUP BY, allows us to apply conditions that **involve aggregate functions**.

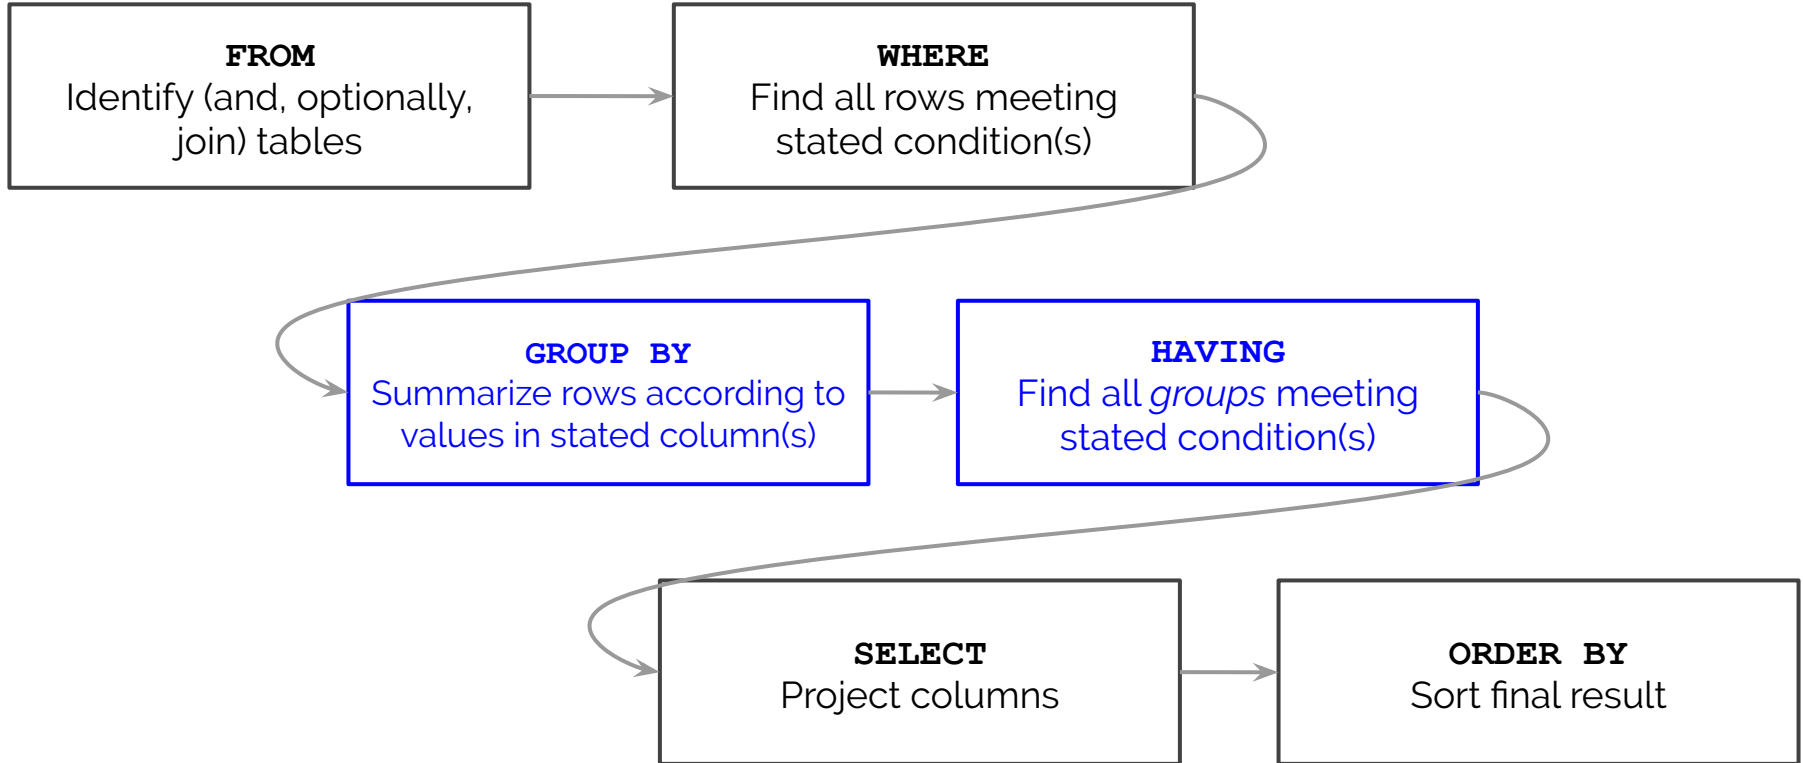
List directors and average gross in million dollars of Movies grouped by director whose average gross is greater than 100 million dollar sorted in the order of the average gross.

```
SELECT director, AVG(gross)/1000000 as grossInMillion  
FROM Movies  
GROUP BY director HAVING grossInMillion > 100  
ORDER BY gorssInMillion;
```

Grouping, Aggregation, and Nulls

- The value NULL is ignored in any aggregation.
 - COUNT(A) where some values in A are NULL counts only non-NULL values in A.
- However, NULL is treated as an ordinary value when forming groups.
- The result of aggregation except for COUNT() over an empty bag of values is NULL. The COUNT of an empty bag is 0.

SQL SELECT Processing Order



Practice Writing Queries

Summary

SQL		Relational Algebra Operator	
SELECT DISTINCT	π	Projection	
FROM	\times	Cartesian product	
INNER JOIN ... ON	\bowtie_{θ}	Theta join	
WHERE	σ	Selection	
AS	ρ	Rename	

Note: This is an oversimplified view. Treat it as a starting place, not the final story.